

Homework 5 Report

Heng Zhang
0029109755
zhan2614

Introduction

In this homework, we use the nn package to build lenet5 as the model to do two tasks. The first task is the recognize number using MNIST dataset. Another task is to recognize object from 100 classes defined in CIFAR100 dataset. The MNIST dataset has 60000 pictures showing numbers from 0 to 9 and each class in the CIFAR100 has 600 images where 500 are training images and 100 are testing images so in total there are 60000 images in CIFAR100.

Setup

I use img2num.py and img2obj.py to implement these two tasks. Both scripts has a cifar100 class. Additionally img2obj.py has view and cam function which enables the functions of viewing the predicted image using opencv and capture real time image from camera and do the online prediction.

Evaluation

Task 1 number recognition

The results for this task are shown in Figures 1, 2, and 3.

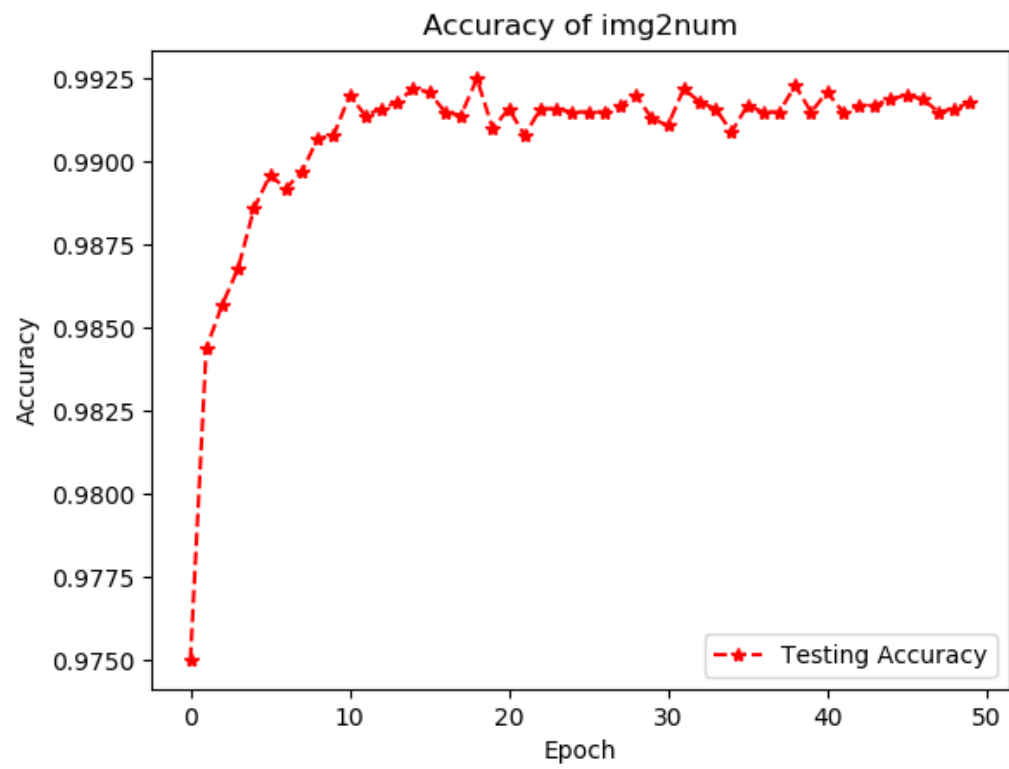


Figure 1 img2num accuracy

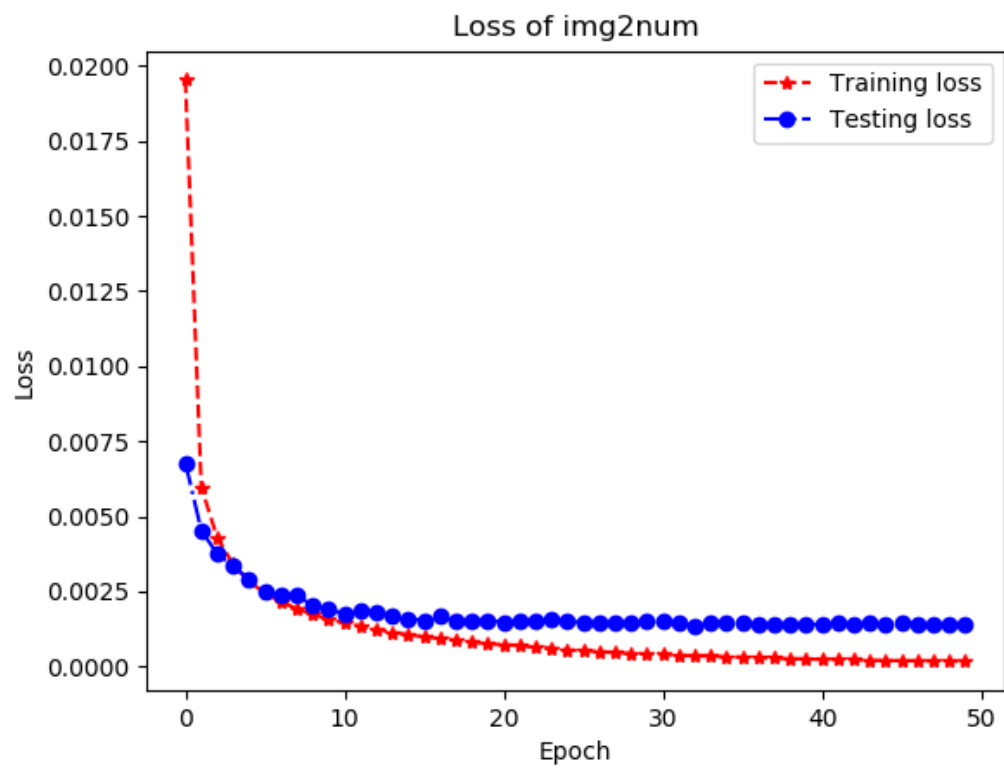


Figure 2 img2num loss

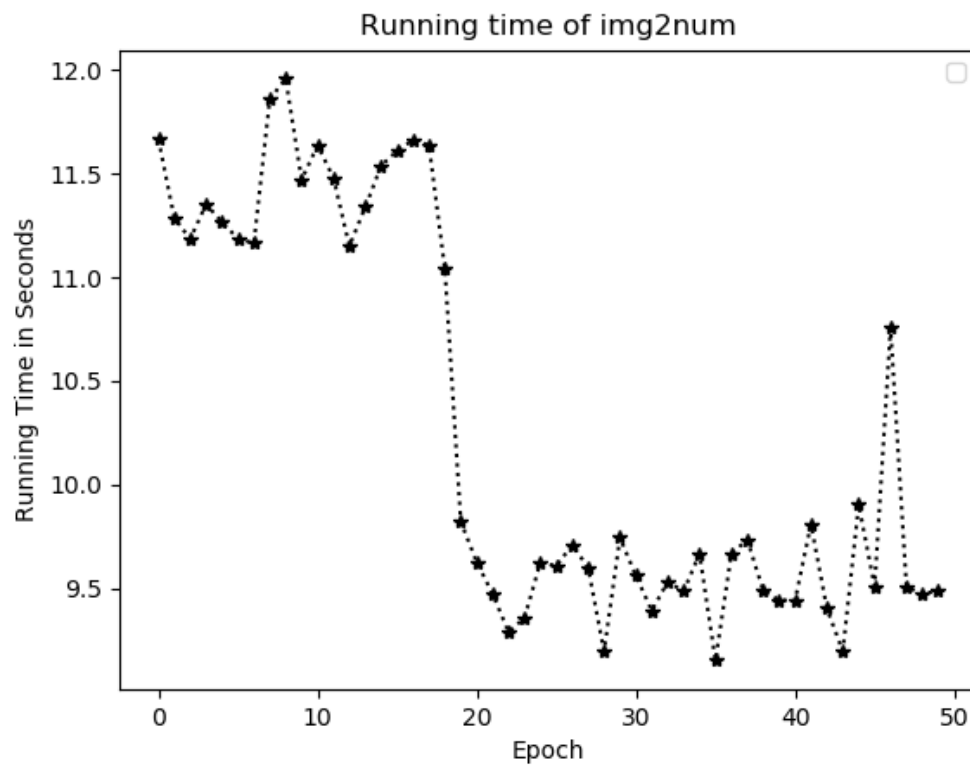


Figure 3 img2num running time

We can see that the accuracy for this task is fairly high to approximately 99%.

Task 2 object observation

The results for this task are represented in Figure 4, 5, and 6.

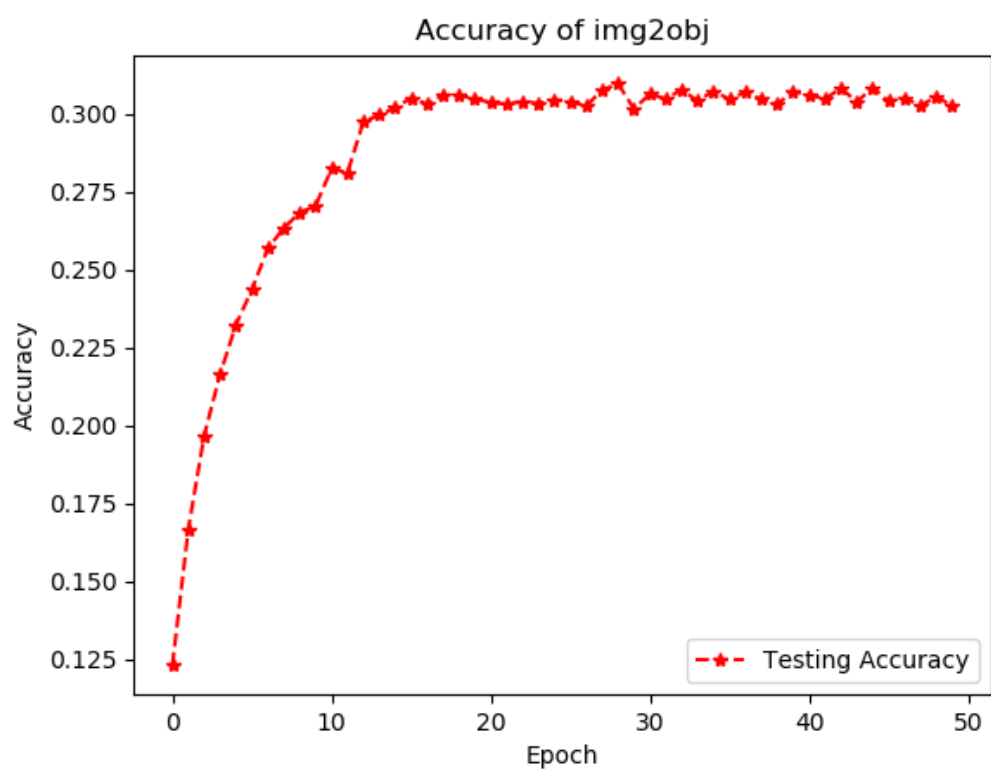


Figure 4 Accuracy of img2obj

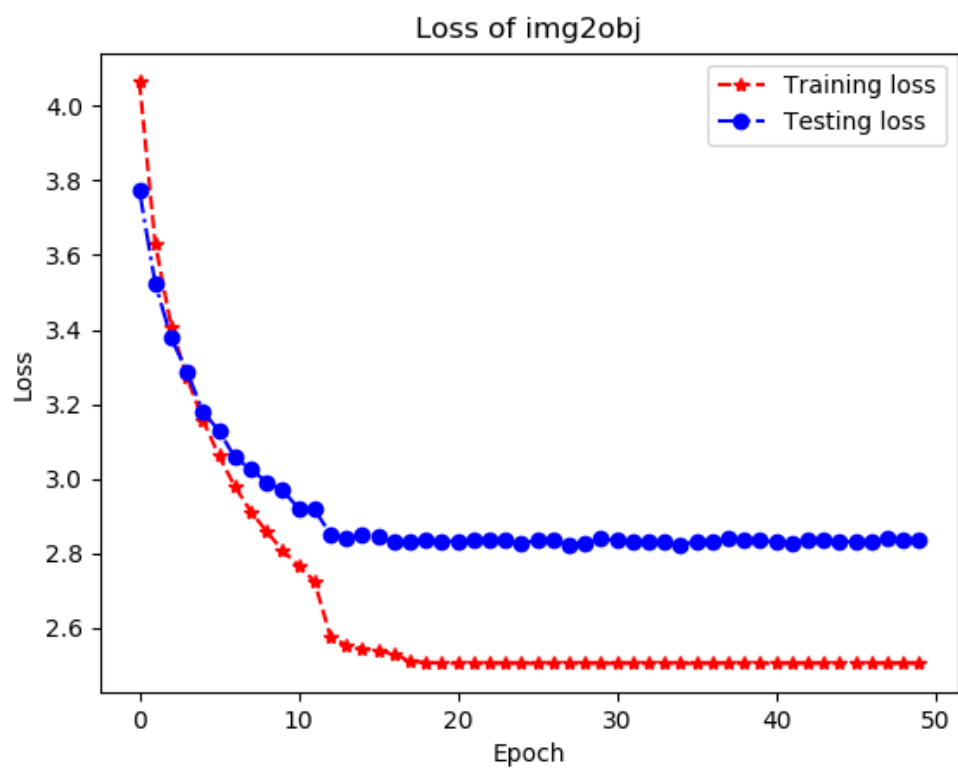


Figure 5 img2obj loss

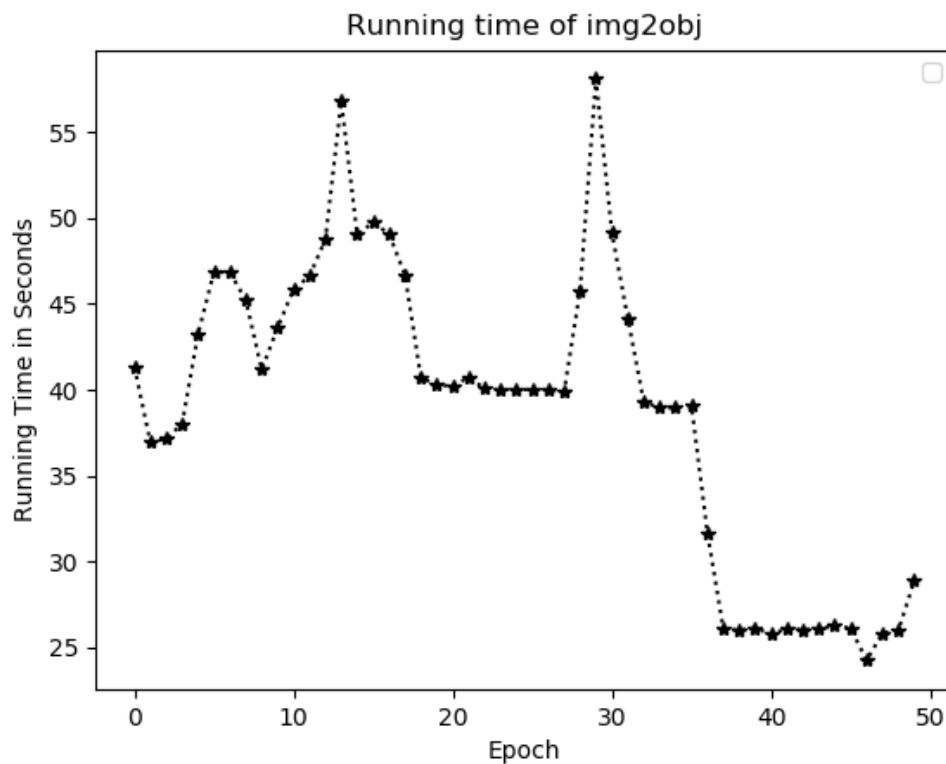


Figure 6 img2obj running time

We see that this task is harder to do and the accuracy approaches about 32%.

In both tasks, the losses are decreasing and the running times are somewhat unstable due to the fact that computers have different available resources on the fly.

Task 3 cam

This task is basically capture a picture from webcam and predicted what is in the picture. Its result is presented in Figure 7. It says the KFC logo and predicts it is a man.

An interesting fact is that sometimes it also predicts it as “cups”, it is also shown in Figure 8.

A finding is that the prediction is not always accurate. In fact most of a time, it is incorrect due to the low prediction accuracy.



Figure 7 opencv capture a man



Figure 8 opencv capture a cup

To regenerate the results in this report, just copy the test.py codes in the Appendix and run it in a machine with opencv and torch installed.

Appendix

img2num.py

```

1. import os
2. import shutil
3. import torch
4. import torch.nn as nn
5. import torch.nn.functional as F
6. from torchvision import datasets, transforms
7. from time import time
8. from torch.autograd import Variable
9.
10.
11. class LeNet(nn.Module):
12.     def __init__(self):
13.         super(LeNet, self).__init__()
14.         self.conv1 = nn.Conv2d(1, 6, 5, padding=2)
15.         self.conv2 = nn.Conv2d(6, 16, 5)
16.         self.fc1 = nn.Linear(5*5*16, 120)
17.         self.fc2 = nn.Linear(120, 84)
18.         self.fc3 = nn.Linear(84, 10)
19.
20.     def forward(self, x):
21.         x = F.relu(self.conv1(x))
22.         x = F.max_pool2d(x, 2)
23.         x = F.relu(self.conv2(x))
24.         x = F.max_pool2d(x, 2)
25.         x = x.view(x.size(0), -1)
26.         x = F.relu(self.fc1(x))
27.         x = F.relu(self.fc2(x))
28.         x = self.fc3(x)
29.         return x
30.
31. class img2num:
32.
33.     def __init__(self):
34.         self.train_batch_size = 60
35.         self.epoch = 50
36.         self.labels = 10
37.         self.rate = 1
38.         self.input_size = 28 * 28
39.         self.test_batch_size = 1000
40.         self.test_loader = torch.utils.data.DataLoader(
41.             datasets.MNIST('./mnist',
42.                             train=False,
43.                             download=True,
44.                             transform=transforms.Compose([transforms.ToTensor()])),

```

```
45.         batch_size=self.test_batch_size, shuffle=True, num_workers=1
46.     0)
47.     self.train_loader = torch.utils.data.DataLoader(
48.         datasets.MNIST('./mnist',
49.             train=True,
50.             download=True,
51.             transform=transforms.Compose([transforms.ToTensor()])),
52.         batch_size=self.train_batch_size, shuffle=True, num_workers=
53.     10)
54.     # input image is 28 * 28 so convert to 1D matrix
55.     # output labels are 10 [0 - 9]
56.     torch.manual_seed(1)
57.     self.model = LeNet()
58.     self.optimizer = torch.optim.SGD(self.model.parameters(), lr=self.ra
59.     te)
60.     self.loss_function = nn.MSELoss()
61.     self.check_point_file = 'img2num_checkpoint.tar'
62.
63.     if os.path.isfile(self.check_point_file):
64.         cp = torch.load(self.check_point_file)
65.         self.start = cp['epoch']
66.         self.best_acc = cp['best_acc']
67.
68.         print('checkpoint found at epoch', self.start)
69.         self.model.load_state_dict(cp['model'])
70.         self.optimizer.load_state_dict(cp['optimizer'])
71.
72.         self.training_loss = cp['training_loss']
73.         self.testing_loss = cp['testing_loss']
74.         self.testing_acc = cp['testing_acc']
75.         self.time = cp['time']
76.     else:
77.         self.start = 0
78.         self.best_acc = 0
79.
80.         self.training_loss = []
81.         self.testing_loss = []
82.         self.testing_acc = []
83.         self.time = []
84.
85.
```



```

86.         # img is 28*28 bytetensor
87.     def forward(self, img):
88.         _3d = torch.unsqueeze(img, 0)
89.         _4d = torch.unsqueeze(_3d, 0)
90.         self.model.eval()
91.         output = self.model(_4d)
92.         _, result = torch.max(output, 1)
93.         return result
94.
95.     def train(self, plot=False):
96.         print('training')
97.         def save(state, better, f=self.check_point_file):
98.             torch.save(state, f)
99.             if better:
100.                 shutil.copyfile(f, 'img2num_best.tar')
101.
102.         def onehot_training(target, batch_size):
103.             output = torch.zeros(batch_size, self.labels)
104.             for i in range(batch_size):
105.                 output[i][int(target[i])] = 1.0
106.             return output
107.
108.         def training():
109.             loss = 0
110.             self.model.train() # set to training mode
111.             for batch_id, (data, target) in enumerate(self.train_loader):
112.                 # data.view change the dimension of input to use forward fu
113.                 forward_pass_output = self.model(data)
114.                 onehot_target = onehot_training(target, self.train_batch_si
115.                 ze)
116.                 #print(onehot_target.type())
117.                 cur_loss = self.loss_function(forward_pass_output, onehot_t
118.                 arget)
119.                 loss += cur_loss.data
120.                 self.optimizer.zero_grad()
121.                 cur_loss.backward()
122.                 self.optimizer.step()
123.                 # loss / number of batches
124.                 avg_loss = loss / (len(self.train_loader.dataset) / self.train
125.                 batch_size)
126.
127.         def testing():

```

```

126.         self.model.eval()
127.         loss = 0
128.         correct = 0
129.         for batch_id, (data, target) in enumerate(self.test_loader):
130.             # data.view change the dimension of input to use forward fu
            nction
131.             forward_pass_output = self.model(data)
132.             onehot_target = onehot_training(target, self.test_batch_siz
            e)
133.             cur_loss = self.loss_function(forward_pass_output, onehot_t
            arget)
134.             loss += cur_loss.data
135.             #print(forward_pass_output.size())
136.             #print(onehot_target.size())
137.             for i in range(self.test_batch_size):
138.                 val, position = torch.max(forward_pass_output.data[i],
                0)
139.                 # print('prediction = {}, actual = {}'.format(int(positi
                on), target[i]))
140.                 if position == target[i]:
141.                     correct += 1
142.             # loss / number of batches
143.             avg_loss = loss / (len(self.test_loader.dataset) / self.test_ba
            tch_size)
144.             accuracy = correct / len(self.test_loader.dataset)
145.             return avg_loss, accuracy
146.         for i in range(self.start + 1, self.epoch + 1):
147.             s = time()
148.             train_loss = training()
149.             e = time()
150.             test_loss, accuracy = testing()
151.             print('Epoch {}, training_loss = {}, testing_loss = {}, accurac
            y = {}, time = {}'.format(i, train_loss, test_loss, accuracy, e - s))
152.             self.testing_acc.append(accuracy)
153.             self.training_loss.append(train_loss)
154.             self.testing_loss.append(test_loss)
155.             self.time.append(e-s)
156.             better = False
157.             if accuracy > self.best_acc:
158.                 better = True
159.             self.best_acc = max(self.best_acc, accuracy)
160.             print('Save checkpoint at', i)
161.             state = {
162.                 'epoch': i,

```

```

163.         'best_acc': self.best_acc,
164.         'model': self.model.state_dict(),
165.         'optimizer': self.optimizer.state_dict(),
166.         'training_loss': self.training_loss,
167.         'testing_loss': self.testing_loss,
168.         'testing_acc': self.testing_acc,
169.         'time': self.time
170.     }
171.     save(state,better)
172.
173.     if plot == True:
174.         return self.time, self.training_loss, self.testing_loss, self.t
esting_acc
175.
176.     '''
177.     plt.plot(range(self.epoch), acc_list, 'r|--', label='Accuracy')
178.     plt.plot(range(self.epoch), train_loss_list, 'b*--
', label='Training Loss')
179.     plt.plot(range(self.epoch), test_loss_list, 'yo--
', label='Test Loss')
180.     plt.xlabel('Epoch')
181.     plt.legend()
182.     plt.title('Library Neural Network Evaluation')
183.     plt.savefig('nn_compare.png')
184.     plt.clf()
185.     '''
186.

```

img2obj.py

```

1. import numpy as np
2. import os
3. import shutil
4. from pprint import pprint as pp
5. import torch
6. import torch.nn as nn
7. import torch.nn.functional as F
8. from torchvision import datasets, transforms
9. from time import time, sleep
10. from torch.autograd import Variable
11. import cv2
12.
13. class LeNet(nn.Module):
14.     def __init__(self):
15.         super(LeNet, self).__init__()

```

```

16.         self.conv1 = nn.Conv2d(3, 6, 5)
17.         self.conv2 = nn.Conv2d(6, 16, 5)
18.         self.fc1 = nn.Linear(5*5*16, 120)
19.         self.fc2 = nn.Linear(120, 84)
20.         self.fc3 = nn.Linear(84, 100)
21.
22.     def forward(self, x):
23.         x = F.relu(self.conv1(x))
24.         x = F.max_pool2d(x, 2)
25.         x = F.relu(self.conv2(x))
26.         x = F.max_pool2d(x, 2)
27.         x = x.view(x.size(0), -1)
28.         x = F.relu(self.fc1(x))
29.         x = F.relu(self.fc2(x))
30.         x = self.fc3(x)
31.         return x
32.
33. class img2obj:
34.
35.     def __init__(self):
36.         self.train_batch_size = 125
37.         self.epoch = 50
38.         self.rate = 0.001
39.         self.input_size = 32 * 32 * 3 #RGB 3 channels of data
40.         self.test_batch_size = 1000
41.         normalize = transforms.Normalize(mean=[0.5, 0.5, 0.5], std=[0.5, 0.5
, 0.5])
42.         self.test_loader = torch.utils.data.DataLoader(
43.             datasets.CIFAR100('./cifar',
44.                               train=False,
45.                               download=True,
46.                               transform=transforms.Compose([transforms.RandomHorizontalFli
p(), transforms.ToTensor(), normalize])),
47.             batch_size=self.test_batch_size, shuffle=True, num_workers=1
0)
48.
49.         self.train_loader = torch.utils.data.DataLoader(
50.             datasets.CIFAR100('./cifar',
51.                               train=True,
52.                               download=True,
53.                               transform=transforms.Compose([transforms.ToTensor(), normali
ze])),
54.             batch_size=self.train_batch_size, shuffle=True, num_workers=
10)

```

```

55.
56.     self.classes = [
57.         'beaver', 'dolphin', 'otter', 'seal', 'whale',
58.         'aquarium fish', 'flatfish', 'ray', 'shark', 'trout',
59.         'orchids', 'poppies', 'roses', 'sunflowers', 'tulips',
60.         'bottles', 'bowls', 'cans', 'cups', 'plates',
61.         'apples', 'mushrooms', 'oranges', 'pears', 'sweet peppers',
62.         'clock', 'computer keyboard', 'lamp', 'telephone', 'television',
63.         'bed', 'chair', 'couch', 'table', 'wardrobe',
64.         'bee', 'beetle', 'butterfly', 'caterpillar', 'cockroach'
65.         ,
66.         'bear', 'leopard', 'lion', 'tiger', 'wolf',
67.         'bridge', 'castle', 'house', 'road', 'skyscraper',
68.         'cloud', 'forest', 'mountain', 'plain', 'sea',
69.         'camel', 'cattle', 'chimpanzee', 'elephant', 'kangaroo',
70.         'fox', 'porcupine', 'possum', 'raccoon', 'skunk',
71.         'crab', 'lobster', 'snail', 'spider', 'worm',
72.         'baby', 'boy', 'girl', 'man', 'woman',
73.         'crocodile', 'dinosaur', 'lizard', 'snake', 'turtle',
74.         'hamster', 'mouse', 'rabbit', 'shrew', 'squirrel',
75.         'maple', 'oak', 'palm', 'pine', 'willow',
76.         'bicycle', 'bus', 'motorcycle', 'pickup truck', 'train',
77.         'lawn-mower', 'rocket', 'streetcar', 'tank', 'tractor'
78.     ]
79.     torch.manual_seed(1)
80.     self.labels = len(self.classes)
81.     # input image is 32 * 32 so convert to 1D matrix
82.     self.model = LeNet()
83.     self.optimizer = torch.optim.Adam(self.model.parameters(), lr=self.rate, weight_decay=0.0005)
84.     self.loss_function = nn.CrossEntropyLoss()
85.     self.check_point_file = 'img2obj_checkpoint.tar'
86.     if os.path.isfile(self.check_point_file):
87.         cp = torch.load(self.check_point_file)
88.         self.start = cp['epoch']
89.         self.best_acc = cp['best_acc']
90.
91.     print('checkpoint found at epoch', self.start)
92.     self.model.load_state_dict(cp['model'])
93.     self.optimizer.load_state_dict(cp['optimizer'])
94.

```

```

95.         self.training_loss = cp['training_loss']
96.         self.testing_loss = cp['testing_loss']
97.         self.testing_acc = cp['testing_acc']
98.         self.time = cp['time']
99.     else:
100.         self.start = 0
101.         self.best_acc = 0
102.
103.         self.training_loss = []
104.         self.testing_loss = []
105.         self.testing_acc = []
106.         self.time = []
107.
108.         # img is 28*28 bytetensor
109.     def forward(self, img):
110.         _4d = torch.unsqueeze(img.type(torch.FloatTensor), 0)
111.         self.model.eval()
112.         output = self.model(_4d)
113.         _, result = torch.max(output, 1)
114.         return self.classes[result]
115.
116.     def train(self, plot=False):
117.         print('training')
118.
119.         def save(state, better, f=self.check_point_file):
120.             torch.save(state, f)
121.             if better:
122.                 shutil.copyfile(f, 'img2obj_best.tar')
123.
124.         def training():
125.             loss = 0
126.             self.model.train() # set to training mode
127.             for batch_id, (data, target) in enumerate(self.train_loader):
128.                 # data.view change the dimension of input to use forward fu
129.                 forward_pass_output = self.model(data)
130.
131.                 cur_loss = self.loss_function(forward_pass_output, target)
132.
133.                 loss += cur_loss.data
134.                 self.optimizer.zero_grad()
135.                 cur_loss.backward()
136.                 self.optimizer.step()
137.             # loss / number of batches

```

```

137.         avg_loss = loss / (len(self.train_loader.dataset) / self.train_
            batch_size)
138.         return avg_loss
139.
140.     def testing():
141.         self.model.eval()
142.         loss = 0
143.         correct = 0
144.         for batch_id, (data, target) in enumerate(self.test_loader):
145.             # data.view change the dimension of input to use forward fu
            nction
146.             forward_pass_output = self.model(data)
147.             cur_loss = self.loss_function(forward_pass_output, target)
148.
149.             loss += cur_loss.data
150.             #print(forward_pass_output.size())
151.             #print(onehot_target.size())
152.             for i in range(self.test_batch_size):
153.                 val, position = torch.max(forward_pass_output.data[i],
                    0)
154.                 # print('prediction = {}, actual = {}'.format(int(position), target[i]))
155.                 if position == target[i]:
156.                     correct += 1
157.             # loss / number of batches
158.             avg_loss = loss / (len(self.test_loader.dataset) / self.test_ba
                tch_size)
159.             accuracy = correct / len(self.test_loader.dataset)
160.             return avg_loss, accuracy
161.
162.         last_acc = 0
163.         for i in range(self.start + 1, self.epoch + 1):
164.             s = time()
165.             train_loss = training()
166.             e = time()
167.             test_loss, accuracy = testing()
168.             print('Epoch {}, training_loss = {}, testing_loss = {}, accurac
                y = {}, time = {}'.format(i, train_loss, test_loss, accuracy, e - s))
169.
170.             if last_acc > accuracy:
171.                 for g in self.optimizer.param_groups:
172.                     g['lr'] = g['lr']/10
173.                     print('learning rate changed to', g['lr'])
174.             last_acc = accuracy
175.             self.testing_acc.append(accuracy)

```

```
174.         self.training_loss.append(train_loss)
175.         self.testing_loss.append(test_loss)
176.         self.time.append(e-s)
177.         better = False
178.         if accuracy > self.best_acc:
179.             better = True
180.         self.best_acc = max(self.best_acc, accuracy)
181.         print('Save checkpoint at', i)
182.         state = {
183.             'epoch': i,
184.             'best_acc': self.best_acc,
185.             'model': self.model.state_dict(),
186.             'optimizer': self.optimizer.state_dict(),
187.             'training_loss': self.training_loss,
188.             'testing_loss': self.testing_loss,
189.             'testing_acc': self.testing_acc,
190.             'time': self.time
191.         }
192.         save(state,better)
193.
194.         #label = self.classes[self.train_loader.dataset[20][1]]
195.         #print("actual label is", label)
196.         #self.view(self.train_loader.dataset[20][0])
197.
198.         if plot == True:
199.             return self.time, self.training_loss, self.testing_loss, self.t
testing_acc
200.
201.     def view(self, img):
202.         category = self.forward(img)
203.         print('Prediction is', category)
204.         img = img.type(torch.FloatTensor) / 2 + 0.5
205.         img_numpy = np.transpose(img.numpy(), (1,2,0))
206.
207.         cv2.namedWindow(category, cv2.WINDOW_NORMAL)
208.         cv2.resizeWindow(category, 640, 480)
209.         cv2.imshow(category, img_numpy)
210.         cv2.waitKey(0)
211.         cv2.destroyAllWindows()
212.
213.     def cam(self, idx = 0):
214.
215.         def prepare(img_origin):
```



```

216.         img_scaled = cv2.resize(img_origin, (32, 32), interpolation=cv2
    .INTER_LINEAR)
217.         # Convert to Tensor and Normalize
218.         prepare = transforms.Compose([transforms.ToTensor(), transforms
    .Normalize(mean=[0.5, 0.5, 0.5], std=[0.5, 0.5, 0.5])])
219.
220.         return prepare(img_scaled)
221.
222.     cam = cv2.VideoCapture(idx)
223.     cam.set(3, 1280)
224.     cam.set(4, 720)
225.     cv2.namedWindow("test")
226.     img_counter = 0
227.     print('Press e/E to exit, c/C to capture a picture\n')
228.     while True:
229.         ret, frame = cam.read()
230.         if not ret:
231.             break
232.         norm_img_tensor = prepare(frame)
233.         predicted_category = self.forward(norm_img_tensor)
234.         print(predicted_category)
235.         cv2.putText(frame, predicted_category, (10,500), cv2.FONT_HERSHEY_
    EY_SIMPLEX, 4, (255, 255, 255), 5, cv2.LINE_AA)
236.         cv2.imshow('Capturing', frame)
237.
238.         k = cv2.waitKey(1) & 0xFF
239.         if k == ord('e'):
240.             # e pressed
241.             print("E hit, closing...")
242.             break
243.         elif k == ord('c'):
244.             # c pressed
245.             img_name = "opencv_frame_{}.png".format(img_counter)
246.             cv2.imwrite(img_name, frame)
247.             print("{} written!".format(img_name))
248.             img_counter += 1
249.
250.     cam.release()
251.
252.     cv2.destroyAllWindows()

```

test.py

```

1. import matplotlib.pyplot as plt
2. from img2num import img2num

```

```

3. from img2obj import img2obj
4. def graph(time, train_loss, test_loss, accuracy, name):
5.
6.     plt.plot(time, 'k*:')
7.     plt.ylabel('Running Time in Seconds')
8.     plt.legend()
9.     plt.xlabel('Epoch')
10.    plt.title("Running time of " + name)
11.    plt.savefig(name+'_time.png')
12.    plt.clf()
13.
14.    plt.plot(range(len(accuracy)), accuracy, 'r*--
        ', label='Testing Accuracy')
15.    plt.xlabel('Epoch')
16.    plt.ylabel('Accuracy')
17.    plt.title("Accuracy of " + name)
18.    plt.legend()
19.    plt.savefig(name + '_acc.png')
20.    plt.clf()
21.
22.    plt.plot(range(len(train_loss)), train_loss, 'r*--
        ', label='Training loss')
23.    plt.plot(range(len(test_loss)), test_loss, 'bo-.', label='Testing loss')
24.    plt.xlabel('Epoch')
25.    plt.legend()
26.    plt.title("Loss of "+name)
27.    plt.ylabel('Loss')
28.    plt.savefig(name+'_loss.png')
29.    plt.clf()
30. print('img2Num testing')
31. img = img2num()
32. time, train_loss, test_loss, accuracy = img.train(True)
33. graph(time, train_loss, test_loss, accuracy, 'img2num')
34. print('img2obj testing')
35. img = img2obj()
36. time, train_loss, test_loss, accuracy = img.train(True)
37. graph(time, train_loss, test_loss, accuracy, 'img2obj')
38. img.cam(0)

```